

5-Stage Pipelined MIPS Processor Optimization

13 Contributions: Pipelining, Forwarding, SIMD, Cache, and Calculus

JUNYI LIU (M143140014)

Computer Architecture, Fall 2025

Prof. Katherine Shu-Min Li

Jan 4, 2026

Outline

1. Project Overview and Motivation
2. Baseline: 5-Stage Pipelined MIPS Architecture
3. Optimization I: Data Forwarding (Contributions 1-4)
4. Optimization II: SIMD with Hardware Unrolling (Contributions 2-3, 5)
5. Optimization III: Branch Prediction (Contribution 6)
6. Advanced Features (Contributions 7-13)
7. Performance Evaluation
8. Conclusion

Project Overview

| Research Scope

Focus on **Parallelism** and **Pipelining** techniques in Verilog HDL:

Technique	Implementation
Pipelining	5-stage datapath (IF/ID/EX/MEM/WB)
ILP	Instruction-level parallelism via pipelining
DLP	Data-level parallelism via SIMD
Hardware Unrolling	Verilog <code>generate</code> construct
Memory Hierarchy	L1 Cache with Write-Back
Calculus Operations	ODE Solver with Euler Method

| My Contributions (1/3)

#	Contribution	Description
1	Data Forwarding	Testbench for CPI measurement + Forwarding
2	Hardware Unrolling	Verilog <code>generate</code> for parallel instantiation
3	SIMD Parallelism	8-lane data-level parallelism demo
4	CPI Analysis	CPI improvement (1.82 to 1.26)
5	SIMD ALU Expansion	Full ALU ops (+, -, x, /, ^)

| My Contributions (2/3)


#	Contribution	Description
6	Branch Prediction	2-bit saturating counter predictor
7	DOE Factorial Analysis	2^4 DOE analysis of optimization effects
8	Expression Parser	Shunting-yard with parentheses
9	CORDIC	Hardware trigonometry (sin/cos)
10	Polynomial Transcendental	exp() and ln() approximation

| My Contributions (3/3)

#	Contribution	Description
11	Cache Hierarchy	L1 Direct-Mapped with Write-Back
12	Integrated Analysis	MiBench + SPEC projection
13	ODE Solver	Forward Euler method (Calculus)

Baseline Architecture

5-Stage Pipelined Datapath



Stage-to-Module Mapping

Stage	Function	Module
IF	Instruction Fetch	IFStage.v
ID	Instruction Decode / Register Read	IDStage.v
EX	ALU Execution	EXStage.v
MEM	Memory Access	MEMStage.v
WB	Register Writeback	WBStage.v

Pipeline Register Implementation

Example: IF/ID Pipeline Register (IF2ID.v)

```
always @ (posedge clk) begin
    if (rst) begin
        PC <= 0;
        instruction <= 0;
    end
    else if (~freeze) begin
        if (flush) begin
            instruction <= 0;
            PC <= 0;
        end
        else begin
            instruction <= instructionIn;
            PC <= PCIn;
        end
    end
end
```

Key Features: Freeze for stalls, Flush for branch mispredictions

Optimization I: Data Forwarding

| Forwarding Mechanism

Design Principle:

Bypass data directly from **EX/MEM stages** to dependent instructions

Multiplexer Selection:

<code>val1_sel</code>	Data Source
<code>2'd0</code>	Normal path
<code>2'd1</code>	Forward from MEM stage
<code>2'd2</code>	Forward from WB stage

Forwarding Implementation

```
module forwarding_EXE (src1_EXE, src2_EXE, ST_src_EXE,  
                      dest_MEM, dest_WB, WB_EN_MEM, WB_EN_WB,  
                      val1_sel, val2_sel, ST_val_sel);  
  
    always @ ( * ) begin  
        {val1_sel, val2_sel, ST_val_sel} <= 0;  
  
        if (WB_EN_MEM && ST_src_EXE == dest_MEM) ST_val_sel <= 2'd1;  
        else if (WB_EN_WB && ST_src_EXE == dest_WB) ST_val_sel <= 2'd2;  
  
        if (WB_EN_MEM && src1_EXE == dest_MEM) val1_sel <= 2'd1;  
        else if (WB_EN_WB && src1_EXE == dest_WB) val1_sel <= 2'd2;  
  
        if (WB_EN_MEM && src2_EXE == dest_MEM) val2_sel <= 2'd1;  
        else if (WB_EN_WB && src2_EXE == dest_WB) val2_sel <= 2'd2;  
    end  
endmodule
```

| Forwarding Results

Configuration	Cycles	Stalls	CPI	IPC
Forwarding OFF	255	114	1.82	0.549
Forwarding ON	182	37	1.26	0.791

Performance Improvements:

- **Stall reduction:** 68% (114 to 37)
- **CPI improvement:** 31% (1.82 to 1.26)
- **IPC increase:** 44% (0.549 to 0.791)

Optimization II: SIMD Unrolling

SIMD Adder Implementation

Module: `simd_demo/simd_add.v`

```
module simd_add #(
    parameter integer LANES = 8,
    parameter integer WIDTH = 8
) (
    input  wire [LANES*WIDTH-1:0] a,
    input  wire [LANES*WIDTH-1:0] b,
    output wire [LANES*WIDTH-1:0] y
);
    genvar i;
    generate
        for (i = 0; i < LANES; i = i + 1) begin : gen_lane
            assign y[i*WIDTH +: WIDTH] =
                a[i*WIDTH +: WIDTH] + b[i*WIDTH +: WIDTH];
        end
    endgenerate
endmodule
```

SIMD ALU Operations


Operation	Code	Example
ADD	3'b000	$a + b$
SUB	3'b001	$a - b$
MUL	3'b010	$a \times b$
DIV	3'b011	a / b
EXP	3'b100	$a ^ b$

Throughput: 8 operations per cycle (8x speedup)

Optimization III: Branch Prediction

| 2-Bit Saturating Counter Predictor

State Machine:



- **Prediction:** Taken if MSB = 1 (states 10 or 11)
- **BHT Size:** 16 entries, indexed by PC[5:2]

Branch Prediction Results

Pattern	Description	Accuracy
Always Taken	Loop branch	90%
Always Not Taken	Conditional	85%
Alternating	Stress test	73%
Loop (9T+1N)	Real loop	78%

Overall Accuracy: 78.33%

Advanced Features

Contribution 7: DOE Analysis

Method: Full Factorial Design ($2^4 = 16$ configurations)

Config	FWD	BP	CACHE	SIMD	Speedup
0	OFF	OFF	OFF	OFF	1.00x
4	OFF	OFF	ON	OFF	2.15x
15	ON	ON	ON	ON	3.94x

Key Finding: Cache provides largest individual improvement

Contribution 8: Expression Parser

Algorithm: Dijkstra's Shunting-yard (1961)

Test Expression	Expected	Result	Status
<code>5 * (3 + 4)</code>	35	35	PASS
<code>2 ^ 3 ^ 2</code>	512	512	PASS
<code>100 / (2 + 3)</code>	20	20	PASS

Features: Parentheses support, right-associativity for $^$

Contribution 9: CORDIC

Algorithm: COordinate Rotation Digital Computer (Volder, 1959)

Specification	Value
Data Width	16-bit fixed-point (Q2.14)
Pipeline Depth	16 stages
Throughput	1 result/cycle
Hardware	Multiplier-free (shift-add only)
Error	Less than 1%

Contribution 10-11: Cache Hierarchy

L1 Cache Parameters:

Parameter	Value
Cache Size	8 KB
Block Size	32 bytes (8 words)
Mapping	Direct-Mapped
Write Policy	Write-Back, Write-Allocate

AMAT Results:

Configuration	AMAT (cycles)	Speedup
No Cache	100.0	1.00x
L1 Only	6.00	16.7x

| Contribution 12: Integrated Analysis

Two Evaluation Methodologies:

Method	Description	Data Source
Actual Execution	Run real program on RTL	MiBench bitcount
Mix-Based Projection	Estimate using instruction mix	SPEC CPU2006

Why two methods?

- Actual execution: Ground truth, limited scope
- Projection: Broader coverage, estimates only

Contribution 13: ODE Solver (Calculus)

Goal: Implement hardware solution for Ordinary Differential Equations

Mathematical Foundation:

State-Space Form: Systems of ODEs can be written as:

$$dX/dt = A * X + B * U$$

Where:

- **X**: State vector (size N)
- **A**: System matrix (N x N)
- **B**: Input matrix (N x M)
- **U**: Input vector (size M)

| Forward Euler Method

Numerical Integration:

$$\begin{aligned} X(t+h) &= X(t) + h * dX/dt \\ &= X(t) + h * (A*X + B*U) \end{aligned}$$

Calculus Operations Implemented:

1. **Differentiation:** Approximate dX/dt as $(X(t+h) - X(t)) / h$
2. **Integration:** Compute $X(t+h) = X(t) + \text{integral of } (dX/dt) dt$
3. **Matrix Operations:** $AX + BU$

ODE Solver vs CORDIC

Aspect	CORDIC (Contrib. 9)	ODE Solver (Contrib. 13)
Algorithm	Shift-add iterations	Matrix mult + Euler
Data Structure	Scalar	Vectors and Matrices
Control Logic	16-stage pipeline	6-state FSM
Memory	None	8K x 64-bit RAM
Math Level	Trigonometry	Calculus

Conclusion: ODE Solver is significantly more complex

ODE Solver Architecture

State Machine Flow:

```
Start -> Prepare -> Interpolate -> Load1 -> Load2 -> Final_Calc
      |         |         |         |         |
      Read N,M  Read h    Calc B*U   Calc A*X   X + h*(A*X+B*U)
```

Modules:

- `Euler.v`: Main solver (279 lines, 6-state FSM)
- `RAM-Pre.v`: 8K x 64-bit dual-port RAM
- `multiplier_16bit.v`: Wallace tree multiplier
- `add_sub_cla.v`: Carry Lookahead Adder

ODE Solver Verified Results

Vivado Command:

```
close_sim -force; set_property top tb_ode_solver [get_filesets sim_1]; launch_simulation; run 5000ns
```

Test Setup:

- $N=2$, $M=1$, $h=1$, A =Identity matrix, $B=[1;1]$, $X_0=[10;20]$, $U=[5]$

Expected Calculation:

$$X(t+h) = [10,20] + 1 * ([10,20] + [5,5]) = [25, 45]$$

Output	Expected	Actual	Status
X_new[0]	25	25	PASS
X_new[1]	45	45	PASS

Performance Evaluation

Overall Performance Summary

Component	Technique	Improvement
Pipeline	5-stage datapath	5x parallelism
Forwarding	Data bypass	31% CPI reduction
SIMD	8-lane parallel	8x throughput
Branch Pred	2-bit counter	78% accuracy
Cache	L1 Direct-Mapped	16.7x AMAT
ODE Solver	Forward Euler	Calculus in HW

| MiBench: Actual Execution Results

Benchmark: MiBench Suite bitcount (Guthaus et al., IEEE 2001)

Configuration	Cycles	Speedup
Baseline (no forwarding)	792	1.00x
With Forwarding	303	2.61x

Key Point: This is actual RTL simulation, not projection.

| SPEC: Mix-Based Projection

Method: Use SPEC CPU2006 instruction mix ratios (Phansalkar et al., ISCA 2007)

Workload	Baseline	Optimized	Speedup
SPEC FP Average	507,000	19,670	25.78x
SPEC INT mcf	649,000	18,626	34.84x
SPEC Overall	533,000	18,028	29.57x

Note: Higher speedup because SPEC has 49.6% memory ops (Cache benefit)

Combined Results

Evaluation Method	Speedup
DOE Analysis (Contribution 7)	3.94x
Cache Hierarchy (Contribution 11)	16.7x (AMAT)
MiBench Actual (Contribution 12)	2.61x
SPEC Projection (Contribution 12)	29.57x

Key Insight: Optimizations are largely orthogonal and combine multiplicatively

Conclusion

| Summary of Achievements

Core Optimizations:

- **Data Forwarding:** 31% CPI improvement (1.82 to 1.26)
- **SIMD Parallelism:** 8x throughput with 8-lane ALU
- **Branch Prediction:** 78.33% accuracy

Advanced Features:

- **Cache Memory:** 16.7x AMAT improvement
- **CORDIC:** Hardware trigonometry (less than 1% error)
- **ODE Solver:** Calculus operations in hardware

Future Work

Technique	Description	Expected Improvement
Superscalar	Dual-issue pipeline	2x IPC increase
Out-of-Order	Dynamic scheduling	Hide latencies
Multi-core	Parallel MIPS cores	2x throughput
Cache Integration	Pipeline + L1 Cache	Real system

Trade-offs: Performance vs Area vs Power consumption

| Key Takeaways

1. Forwarding as Critical Optimization

- Empirical data shows 31% CPI improvement

2. Cache Provides Largest Benefit

- 16.7x AMAT improvement addresses Memory Wall

3. Hardware Calculus is Feasible

- ODE Solver demonstrates differential equations in Verilog

4. Power of Verilog `generate`

- Enables parametric hardware unrolling for parallelism